



RopGun

Detecting ROP attacks using modern HW facilities

by anticlockwise



Let's begin



DEMO



Outline



1. What is a ROP exploit?
2. How do Modern Processors work?
3. Let's use HW to save the world from ROP!



Outline




1. What is a ROP exploit?



Buffer overflow

```
void foo() {  
  
    char buffer[10];  
  
    scanf("%s", buffer); //DOH  
  
}
```






Spawn a Shell!



Let's override the return address content!

After foo() the execution will be resumed from the return address.






Spawn a Shell!



For example:

“aaaaaaaaaa” + “aaaa” + addr_to_return



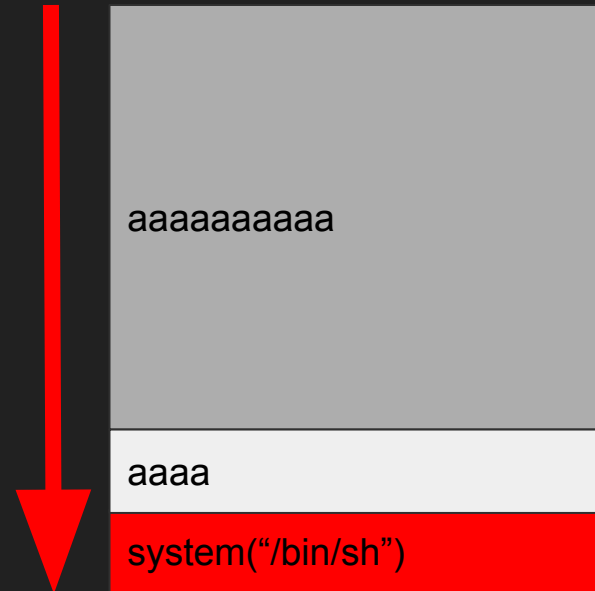


Spawn a Shell!

What to put in `addr_to_return`?

Calling `system("/bin/sh")` will execute a shell in the current context

But what if we don't have `system` available to call?





Spawn a Shell!



Let's dissect system!

Browse the source code of [glibc/sysdeps/posix/system.c](#)
do_system

```
116 #ifdef FORK
117     pid = FORK ();
118 #else
119     pid = __fork ();
120 #endif
121     if (pid == (pid_t) 0)
122     {
123         /* Child side. */
124         const char *new_argv[4];
125         new_argv[0] = SHELL_NAME;
126         new_argv[1] = "-c";
127         new_argv[2] = line;
128         new_argv[3] = NULL;
129
130         /* Restore the signals. */
131         (void) __sigaction (SIGINT, &intr, (struct sigaction *) NULL);
132         (void) __sigaction (SIGQUIT, &quit, (struct sigaction *) NULL);
133         (void) __sigprocmask (SIG_SETMASK, &omask, (sigset_t *) NULL);
134         INIT_LOCK ();
135
136         /* Exec the shell. */
137         (void) __execve (SHELL_PATH, (char *const *) new_argv, __environ);
138         _exit (127);
139     }
```



Spawn a Shell!



__execve is the stub for the execve system call: the system call that transforms the calling process into a new process to execute.

It sets:

eax -> 0x0b

ebx -> address of “/bin/sh” (line to execute)

ecx -> NULL

edx -> NULL

then executes:

int \$0x80



Spawn a Shell!



We need to execute something similar to:

```
mov    ebx, /bin/sh_string_address
mov    ecx, 0x0
mov    edx, 0x0
mov    eax, 0xb
int    0x80
```

But how?


... Let's notice something



Composing a chain

If `addr_to_return` point to some code that ends in “ret”

the execution will then continue from the address after



...and we can iterate again and again



Composing a chain




So why we don't insert more than one address to return to?

For example:

“aaaaaaaaaa” + “aaaa” + fn1_addr + fn2_addr

↑ ↑
padding (10) ebp





Composing a chain

fn1:


```
mov eax, 0xabadcafe
```

```
ret
```

fn2:

```
mov ecx, eax
```

```
ret
```





Composing a chain

foo:


...

ret ← EIP

EAX: ?

ECX: ?

ESP →





Composing a chain



fn1:


mov eax, 0xabadcafe ← EIP

ret

EAX: ?

ECX: ?

ESP →





Composing a chain

fn1:


```
mov eax, 0xabadcafe
```

```
ret ← EIP
```

EAX: 0xabadcafe

ECX: ?

ESP →





Composing a chain

fn2:


mov ecx, eax ← EIP

ret

EAX: 0xabadcafe

ECX: ?

ESP →





Composing a chain

fn2:


```
mov ecx, eax
```

```
ret ← EIP
```

EAX: 0xabadcafe

ECX: 0xabadcafe

ESP →





Composing a chain



We can insert addresses to return to execute, as long as we want!

BUT: every piece of code we execute must end with a ret instruction, to continue the chain

Return address (1st function)

Return address (2nd function)

Return address (3rd function)

Return address (4th function)

Return address (5th function)



Composing a chain



For example, let's load ebx with the address of a /bin/sh string

```
mov ebx, writable_address; ret;
```

```
pop eax; ret;
```

```
"/bin"
```

```
mov [ebx], eax; ret;
```

```
pop eax; ret;
```

```
"/sh"
```

```
mov [ebx+4], eax; ret;
```

Now ebx points to the /bin/sh string



Outline




1. What is a ROP exploit?
2. How do Modern Processors work?



AMD Ryzen

The
Roman
Xploit



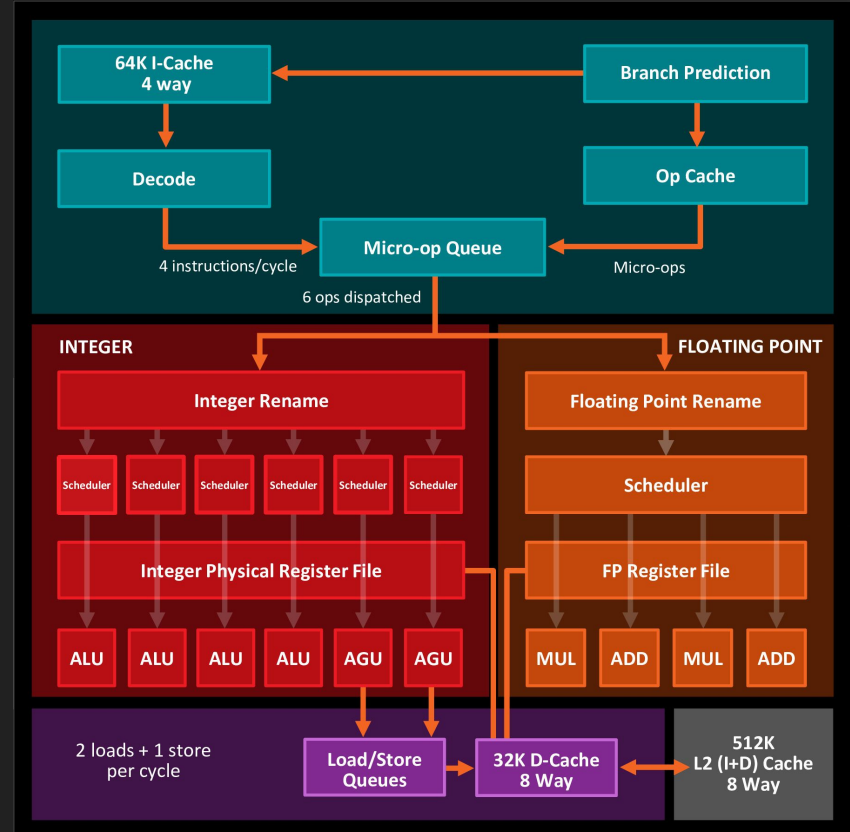


AMD Ryzen





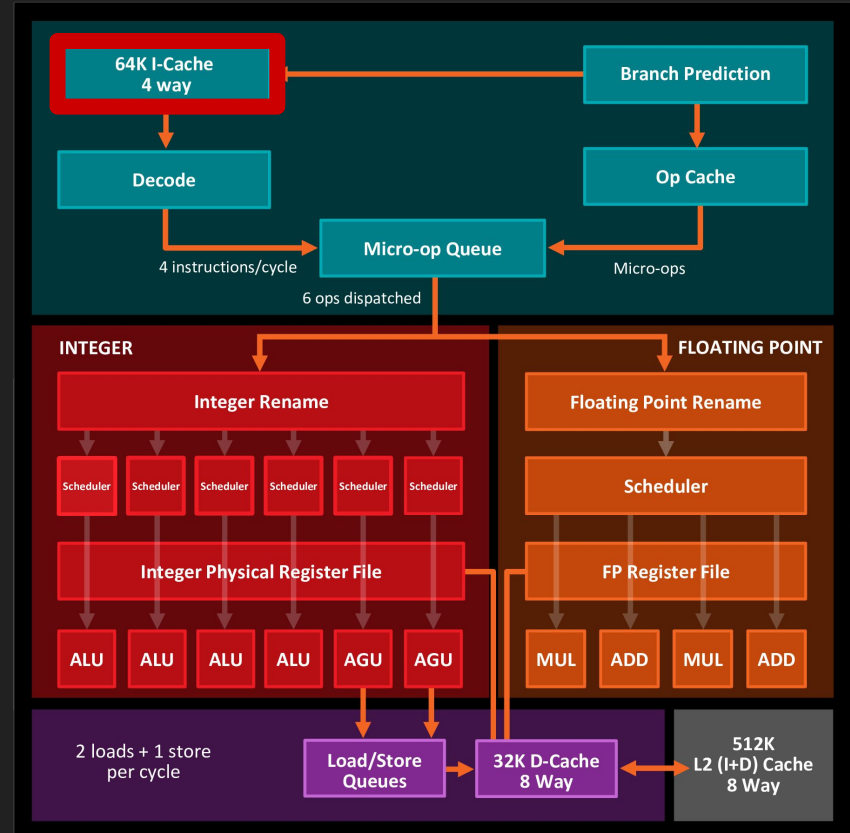
AMD Ryzen





AMD Ryzen

add qword ptr [rax], rbx

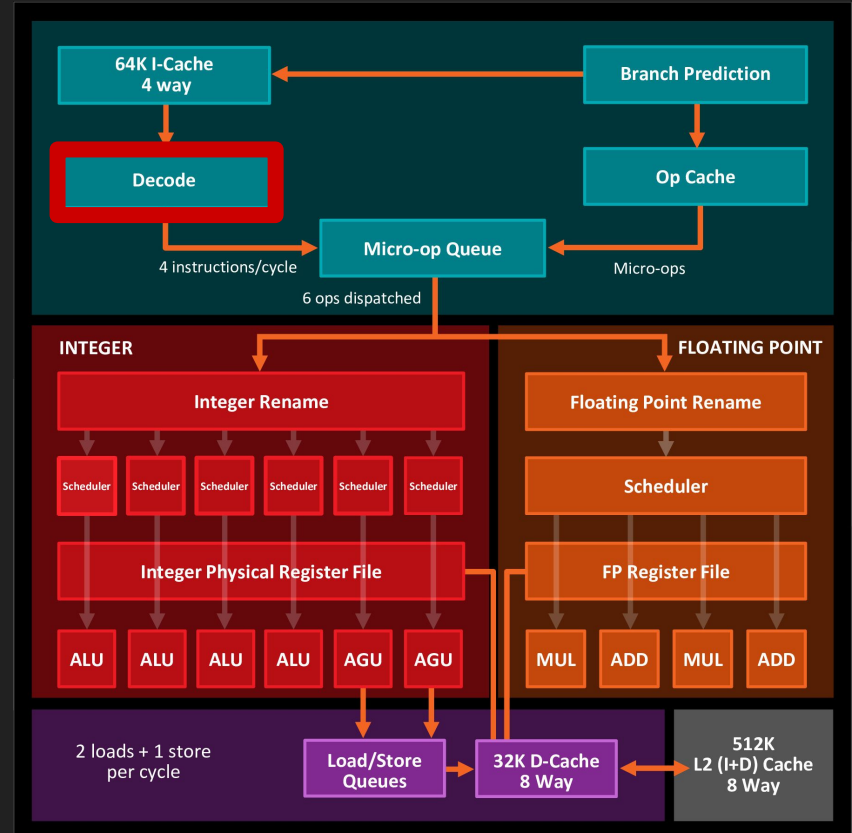




AMD Ryzen



add qword ptr [rax], rbx

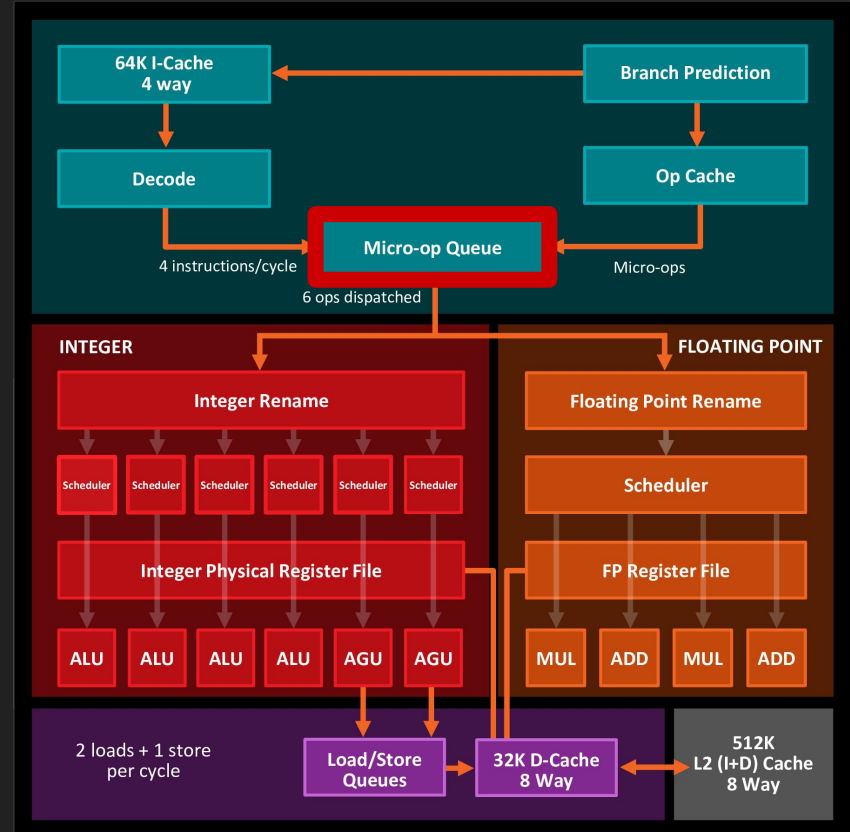





AMD Ryzen



add qword ptr [rax], rbx

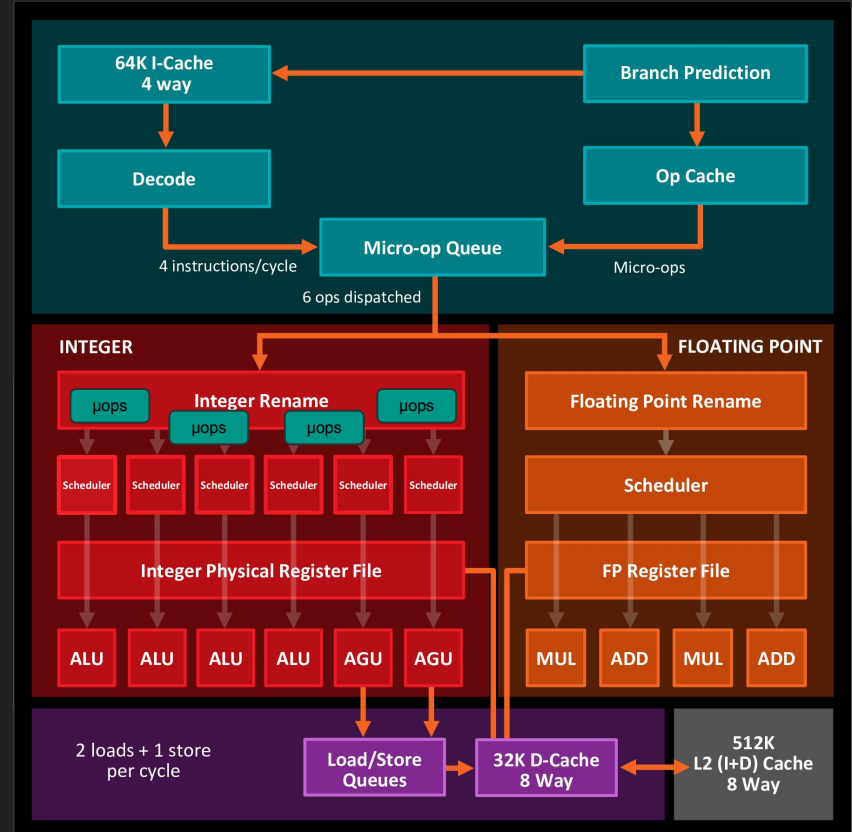





AMD Ryzen



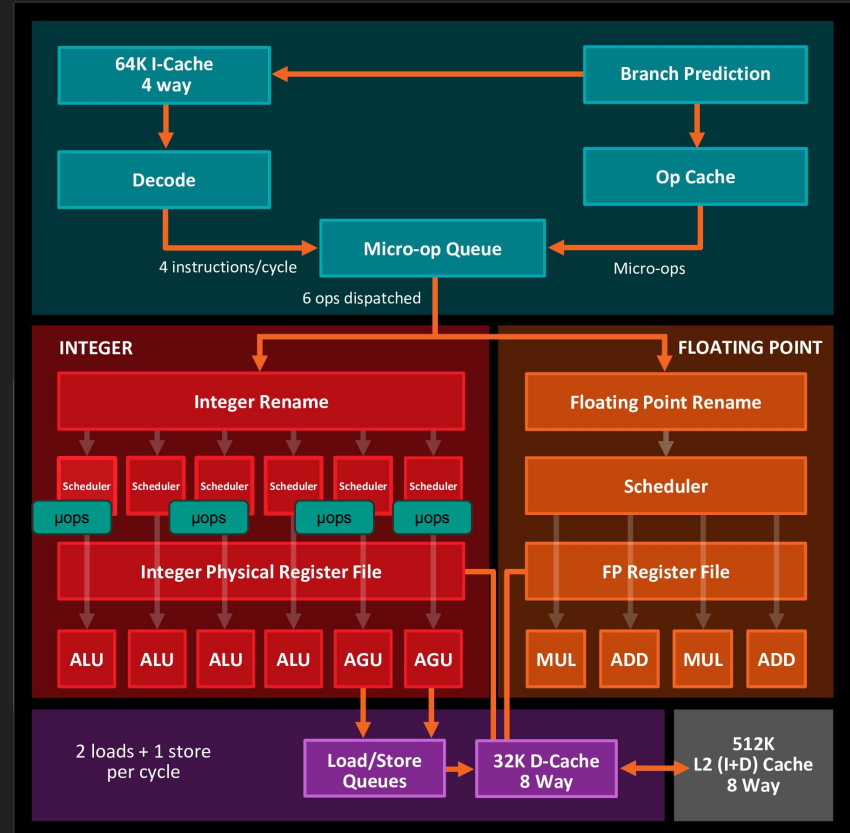

add qword ptr [rax], rbx





AMD Ryzen

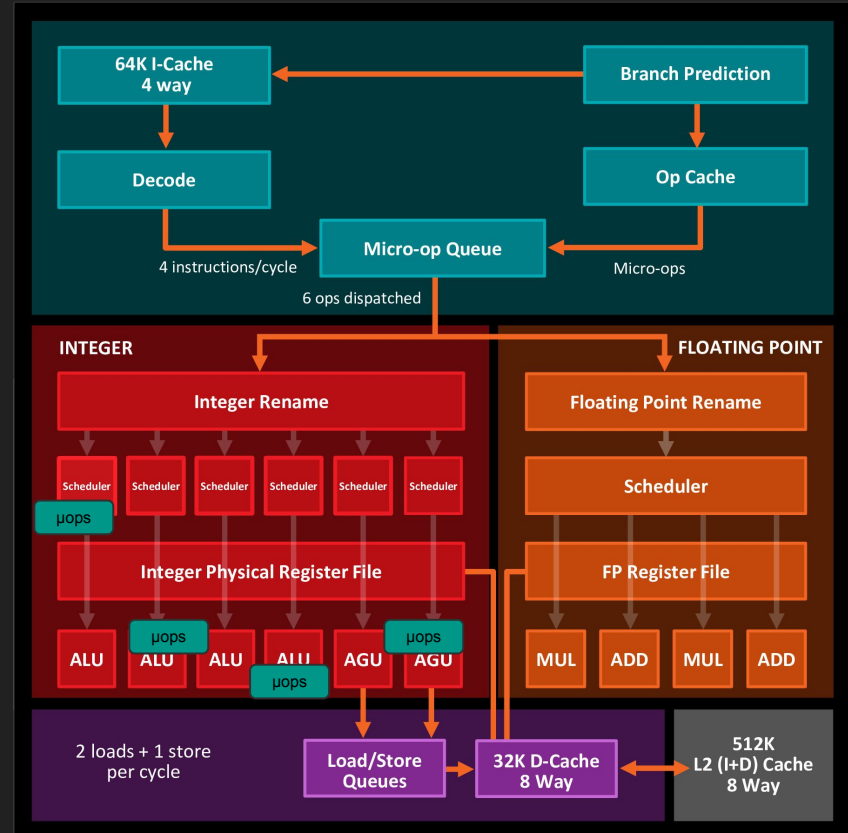

add qword ptr [rax], rbx





AMD Ryzen

add qword ptr [rax], rbx






AMD Ryzen

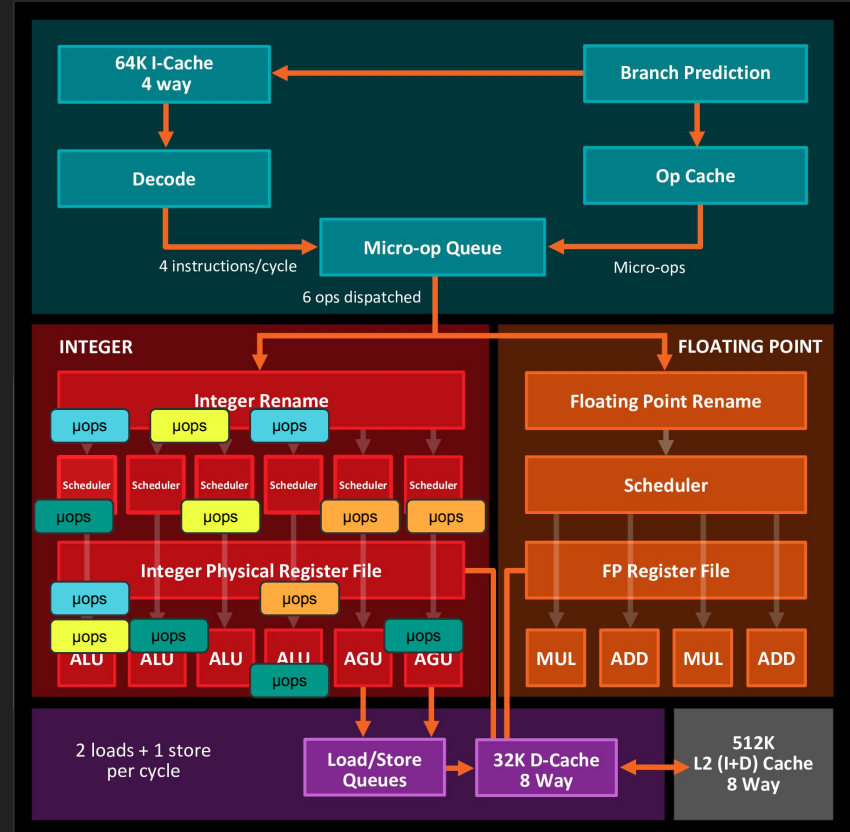


add qword ptr [rax], rbx



mov rdx, 1

cmp rdx, qword ptr [rax]

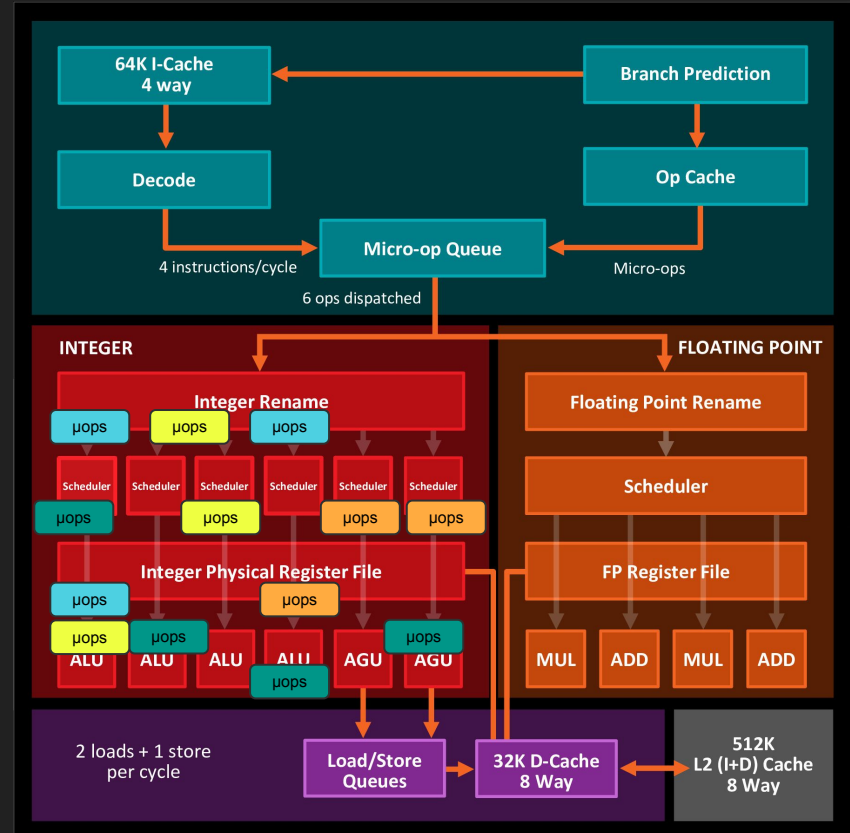
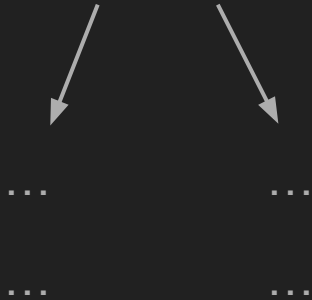




AMD Ryzen



```
add qword ptr [rax], rbx
mov rdx, 1
cmp rdx, qword ptr [rax]
jne 0xdeadbeef
```

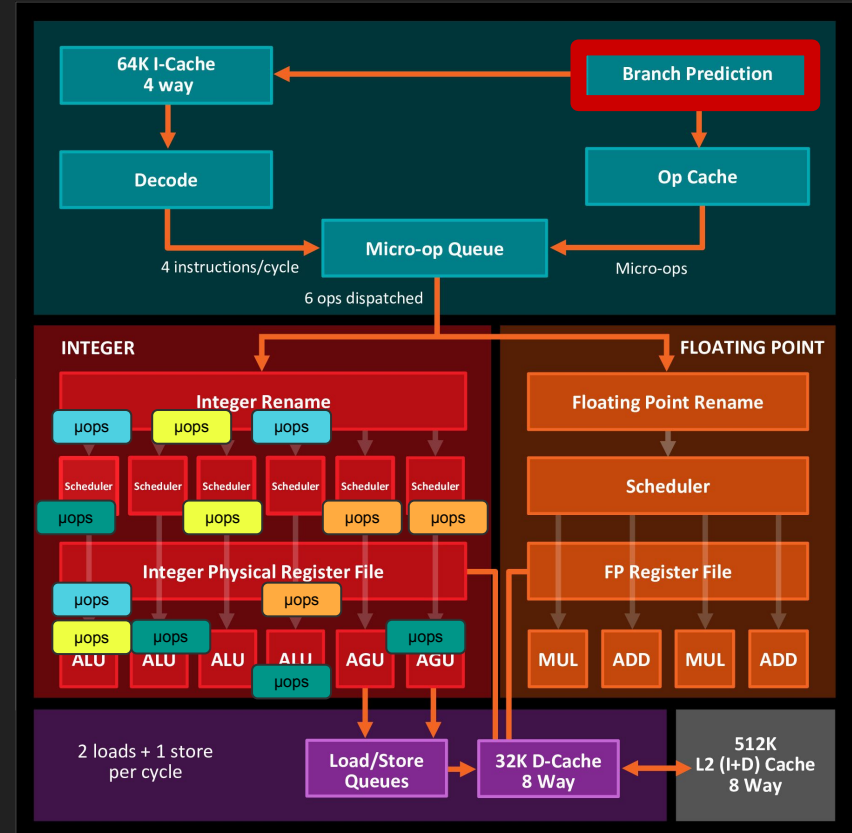





AMD Ryzen



```
add qword ptr [rax], rbx
mov rdx, 1
cmp rdx, qword ptr [rax]
jne 0xdeadbeef
```





Branch Prediction



```
void foo(int i) {  
    int a = 3;  
    int b = 5;  
    if(i > 0) ←  
        bar(a, b); ←  
}
```

```
int bar(int x, int y) {  
    return x+y; ←  
}
```



Branch Prediction



```
void foo(int i) {  
    int a = 3;  
    int b = 5;  
    if(i > 0)  
        bar(a, b);  
}
```


```
int bar(int x, int y) {  
    return x+y;  
}
```





Return Stack Buffer

```
void foo(int i) {  
    int a = 3;  
    int b = 5;  
    if(i > 0)  
        bar(a, b);  
}  
  
int bar(int x, int y) {  
    return x+y;  
}
```





Return Stack Buffer

```
void foo(int i) {  
    int a = 3;  
    int b = 5;  
    if(i > 0)  
        bar(a, b);  
}  
  
int bar(int x, int y) {  
    return x+y;   
}
```



Return Address
Parameters
Local Variables
Saved RBP
Previous Return Address





Return Stack Buffer



```
void foo(int i) {  
    int a = 3;  
    int b = 5;  
    if(i > 0)  
        bar(a, b);  
}
```

```
int bar(int x, int y) {  
    return x+y;   
}
```







Return Stack Buffer

```
void foo(int i) {  
    int a = 3;  
    int b = 5;  
    if(i > 0)  
        bar(a, b);  
}
```

```
int bar(int x, int y) {  
    return x+y;  
}
```



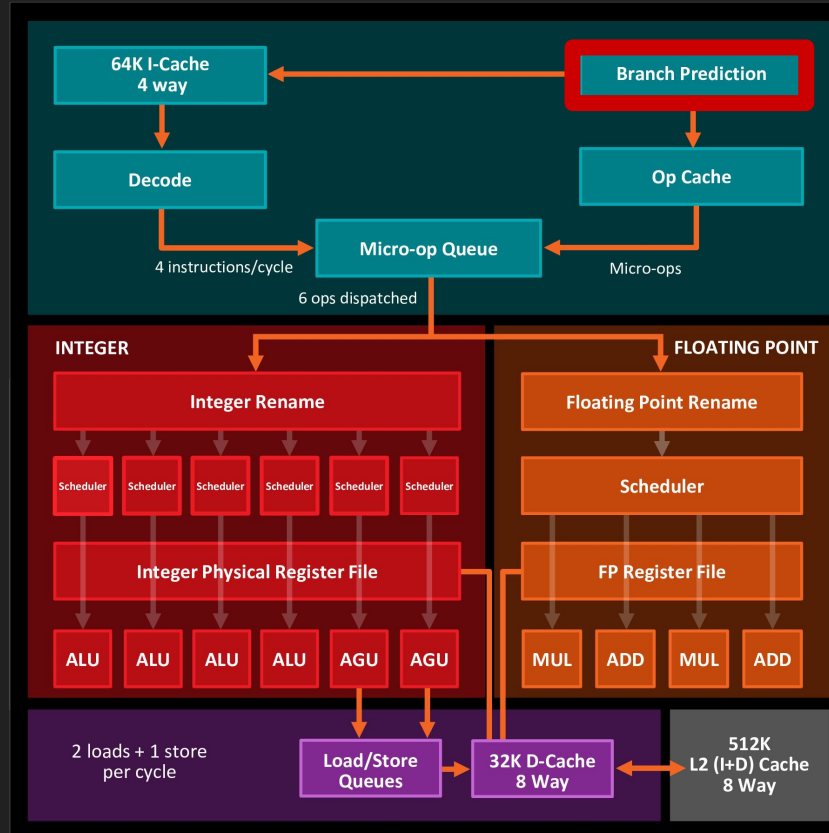


Return Stack Buffer

ret



?



Return Address
...
...
...



Stack



Outline



1. What is a ROP exploit?
2. How do Modern Processors work?
3. Let's use HW to save the world from ROP!



Let's notice something...



We will write our ROPchain on the stack

It will execute a lot of ret

But no call has ever inserted any right address in the RAS

Return Address
...
...

buffer
Saved RBP
Evil ret address
Evil ret address
Evil ret address
Evil ret address
Evil ret address



Let's notice something...

Return Address
...
...
...
...
...
...

← wrong prediction



buffer
Saved RBP
Evil ret address
Evil ret address
Evil ret address
Evil ret address
Evil ret address



Let's notice something...

Return Address
...
...
...
...
...
...

← wrong prediction



buffer
Saved RBP
Evil ret address
Evil ret address
Evil ret address
Evil ret address
Evil ret address



Let's notice something...

Return Address
...
...
...
...
...
...

← wrong prediction



buffer
Saved RBP
Evil ret address
Evil ret address
Evil ret address
Evil ret address
Evil ret address



Let's notice something...

Return Address
...
...
...
...
...
...

← wrong prediction

High rate of misprediction in
the system!

buffer
Saved RBP
Evil ret address
Evil ret address
Evil ret address
Evil ret address
Evil ret address





RopGun



Now the idea is simple!

Just monitor the misprediction rate (through performance counters) in the system to detect ROP payloads executing!

Kill a process when the misprediction rate is too high



How?

Use Performance Monitoring Counters!

```
> perf stat -e branches:u,branch-misses:u -B ./fib 50000 1
500000

Performance counter stats for './fib 50000 1':

      1.830.105      branches:u
           3.266      branch-misses:u          #      0,18% of all branches

0,002706521 seconds time elapsed
```



RopGun

The
Roman
Xploit

DEMO 2



Thank you!



Questions?